

AI-Code_Reviewer Report and fixes

Captured, 8 October 2026. The ai-code-reviewer agent (.claude/agents/ai-code-reviewer.md) run on branch tutor-model-label, git diff main...HEAD at commit 7ee9455, before merging.

What the pull request does

- The Tutor page names the model answering: the OpenRouter slug, made readable by modelLabel in src/lib/tutor.ts ("Model: Claude Sonnet 5.5 by Anthropic", the slug itself on hover). The server page reads the slug with tutorModel() and passes only the name and slug to the client component.
- CLAUDE.md describes the app as built around the AI grammar tutor and why it needs the AI, and gains the "AI model calls" and "Embeddings" rules, including why the embeddings use Qwen3. README.md mentions the model line.

Verdict

Nothing blocking. No secret can reach the browser through this change, and the documentation's claims match the code.

- tutorServer.ts starts with import "server-only" and is imported only by the server page; TutorChat.tsx imports nothing from it and reads no environment variables.
- modelLabel lives in tutor.ts, which reads no environment variables. Only { label, slug } is sent to the client; nothing derived from the key.
- The page shows the same model that answers: tutorModel() is what the tutor and merge routes send to OpenRouter.
- The only OpenRouter calls are in tutorServer.ts; the embedding section matches the comment in tutor.ts and the migration's halfvec(1024) column.

Findings and what was done

Severity	Finding	What was done
Warning	modelLabel got some real slugs wrong: "Deepseek Chat", "GLM 4.6 by Z Ai", "Kimi K2 by Moonshotai", "O3 Mini", "Auto by Openrouter"; a test locked in "Deepseek". The reviewer's smallest fix was to show the raw slug instead.	Kept readable names, since that was the request, and fixed them: one spelling table for makers and model words alike (DeepSeek, Z.ai, Moonshot AI, OpenRouter, GPT, GLM, OSS), OpenAI's o-series kept in lower case. Tests cover each case, plus an unknown maker.
Suggestion	The heuristic will fall behind OpenRouter's catalogue; note that unknown makers are only given a capital.	The table's comment now says so; an unlisted maker reads plainly ("Some Lab"), never wrongly guessed.
Suggestion	A test named "returns a slug it cannot read as it is" checks a trimmed result.	Renamed to say it is returned unchanged but trimmed.

Suggestion	"Model: ..." showed even where no model would answer: no studied language, trial used, daily limit, storage full.	Shown only where a question can be asked; a new test checks it is absent with no language and at the daily limit.
Suggestion	CLAUDE.md repeated the model slug, embedding model and dimensions in two or three places.	The Stack table's Tutor row now points at the new sections instead of repeating them.
Suggestion	"Never change the embedding model" was followed straight away by how to change it.	Reworded: if it ever must change, it cannot be done without the migration that clears every stored embedding.

Not acted on: when the OpenRouter key is not configured the page still names the model. The browser cannot know that, and the tutor already says it failed when asked; the reviewer listed it as fine if intended.

Checks after the fixes

Check	Result
Type check, lint (npx tsc --noEmit, npx eslint src/)	Pass
Unit tests (npx vitest run)	851 of 851, 5 new
Production build (npm run build)	Pass; /tutor still server-rendered

Commits

- 7ee9455 Name the tutor's model on the Tutor page; describe the AI in CLAUDE.md
- 433bbda Apply the code reviewer's findings to the model label